

```
#import all needed packages up front
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_validate
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
import matplotlib.pyplot as plt
import seaborn as sns
```

It is very important for companies to understand why and when they will lose employees. This view of attrition can aid them in retaining top talent while also better planning for loss.

Here we have a sample dataset where key features of employees within a company have been kept and a value for whether that employee left (Yes or No) was provided at a snapshot in time.

Our goal here will be to clean the data, select and tune a feature set, then create and compare a few models which could assist the company in question address the attrition problem.

Credit goes to Prashant Patel (<https://www.kaggle.com/patelprashant>) from where the dataset has been collected

```
#read in the dataset
df=pd.read_csv("https://www.dropbox.com/s/cl/fi/zkaa4dkvip5tkuhdtrrru/employee_attrition.csv?rlkey=urstk5jdsezbzj1da69bvjgff1&st=czxxh2kz&d1=1")
df.head()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	E
0	50.0	No	Travel_Rarely	1126.0	Research & Development	1.0	2	Medical	1	997	
1	36.0	No	Travel_Rarely	216.0	Research & Development	6.0	2	Medical	1	178	
2	21.0	Yes	Travel_Rarely	337.0	Sales	7.0	1	Marketing	1	1780	
3	50.0	No	Travel_Frequently	1246.0	Human Resources	NaN	3	Medical	1	644	
4	52.0	No	Travel_Rarely	994.0	Research & Development	7.0	4	Life Sciences	1	1118	

```
#adjust the target variable to be 1 or 0
df["Attrition"]=df["Attrition"].apply(lambda x: 1 if x=="Yes" else 0)
```

One quick view of what our dataframe looks like.

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1029 entries, 0 to 1028
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    893 non-null    float64
1   Attrition                             1029 non-null   int64
2   BusinessTravel                        1024 non-null   object
3   DailyRate                             1002 non-null   float64
4   Department                            1029 non-null   object
5   DistanceFromHome                      934 non-null    float64
6   Education                             1029 non-null   int64
7   EducationField                        1029 non-null   object
8   EmployeeCount                         1029 non-null   int64
9   EmployeeNumber                        1029 non-null   int64
10  EnvironmentSatisfaction               1029 non-null   int64
11  Gender                                1029 non-null   object
12  HourlyRate                            1029 non-null   int64
13  JobInvolvement                        1029 non-null   int64
14  JobLevel                              1029 non-null   int64
15  JobRole                               1029 non-null   object
```

```

16 JobSatisfaction          1029 non-null int64
17 MaritalStatus            1024 non-null object
18 MonthlyIncome            1029 non-null int64
19 MonthlyRate              1029 non-null int64
20 NumCompaniesWorked       1029 non-null int64
21 Over18                   1029 non-null object
22 OverTime                  1029 non-null object
23 PercentSalaryHike        1029 non-null int64
24 PerformanceRating        1029 non-null int64
25 RelationshipSatisfaction  1029 non-null int64
26 StandardHours            1029 non-null int64
27 StockOptionLevel         1029 non-null int64
28 TotalWorkingYears        1029 non-null int64
29 TrainingTimesLastYear    1029 non-null int64
30 WorkLifeBalance          1029 non-null int64
31 YearsAtCompany           1029 non-null int64
32 YearsInCurrentRole       1029 non-null int64
33 YearsSinceLastPromotion  1029 non-null int64
34 YearsWithCurrManager     1029 non-null int64
dtypes: float64(3), int64(24), object(8)
memory usage: 281.5+ KB

```

We want to take a look at the object columns as these will need to be converted so that they are represented numerically.

```

for col in df.select_dtypes(include='object').columns:
    print(df[col].value_counts())

```

```

BusinessTravel
Travel_Rarely      723
Travel_Frequently  199
Non-Travel         102
Name: count, dtype: int64
Department
Research & Development  676
Sales                  311
Human Resources         42
Name: count, dtype: int64
EducationField
Life Sciences      426
Medical            328
Marketing          110
Technical Degree   82
Other              66
Human Resources    17
Name: count, dtype: int64
Gender
Male      617
Female    412
Name: count, dtype: int64
JobRole
Sales Executive      217
Research Scientist   214
Laboratory Technician 179
Manufacturing Director 95
Healthcare Representative 89
Manager              73
Sales Representative  66
Research Director    62
Human Resources      34
Name: count, dtype: int64
MaritalStatus
Married      474
Single       320
Divorced     230
Name: count, dtype: int64
Over18
Y      1029
Name: count, dtype: int64
OverTime
No      731
Yes     298
Name: count, dtype: int64

```

We are going to drop some columns which are missing a large number values, have no variability, and/or which clearly have collinearity. For the remaining object columns, we will use `get_dummies` for those, while ensuring we drop one from each set.

```

#use get_dummies to convert object columns to numeric features
df=pd.get_dummies(df, columns=['Gender','OverTime','BusinessTravel','Department','EducationField','JobRole','MaritalStatus'], drop_first=True)

```

```
#drop columns that will not be useful, and may be harmful, to our model
df.drop(columns=["Over18", "DistanceFromHome", "EmployeeCount", "EmployeeNumber", "DailyRate", "HourlyRate"], inplace=True)
```

From our initial view of the data, we saw that Age was missing a large number of values. However, where attrition is concerned, it is reasonable to assume that Age could be a significant feature. Therefore, we will create a very simple linear regression model to fill in the missing Age values. We will select features for that very simple model based on correlations to the Age field.

```
#look at all correlations versus the Age field
correlation_matrix = df.corr()
age_correlations = correlation_matrix['Age']
print(age_correlations)
```

```
Age 1.000000
Attrition -0.175555
Education 0.226646
EnvironmentSatisfaction -0.024051
JobInvolvement 0.024176
JobLevel 0.506348
JobSatisfaction 0.007670
MonthlyIncome 0.492360
MonthlyRate 0.018697
NumCompaniesWorked 0.298354
PercentSalaryHike 0.011593
PerformanceRating 0.007457
RelationshipSatisfaction 0.062533
StandardHours NaN
StockOptionLevel 0.023282
TotalWorkingYears 0.676650
TrainingTimesLastYear -0.026077
WorkLifeBalance -0.042949
YearsAtCompany 0.291798
YearsInCurrentRole 0.218028
YearsSinceLastPromotion 0.178679
YearsWithCurrManager 0.186070
Gender_Male 0.012663
OverTime_Yes -0.013611
BusinessTravel_Travel_Frequently -0.049995
BusinessTravel_Travel_Rarely 0.072931
Department_Research & Development 0.006722
Department_Sales -0.026030
EducationField_Life Sciences -0.014555
EducationField_Marketing 0.037174
EducationField_Medical 0.016565
EducationField_Other -0.044943
EducationField_Technical Degree -0.006064
JobRole_Human Resources -0.025356
JobRole_Laboratory Technician -0.139663
JobRole_Manager 0.285131
JobRole_Manufacturing Director 0.037921
JobRole_Research Director 0.187800
JobRole_Research Scientist -0.158640
JobRole_Sales Executive 0.026403
JobRole_Sales Representative -0.208113
MaritalStatus_Married 0.113787
MaritalStatus_Single -0.132596
Name: Age, dtype: float64
```

```
#create subsets of rows with and without Age value
hasAge=df[df['Age'].notna()]
missingAge=df[df['Age'].isna()]
```

```
#based on the correlation matrix, define features for our linear regression model
features=['TotalWorkingYears', 'JobLevel', 'MonthlyIncome']
```

```
#fit the model
reg=LinearRegression()
reg.fit(hasAge[features], hasAge['Age'])
```

```
#predict and assign missing Age values
df.loc[df['Age'].isna(), 'Age']=reg.predict(missingAge[features])
```

To further eliminate non-useful features, we will programmatically drop any which have no variance.

```
zero_variance=df.columns[df.var()== 0]
df.drop(columns=zero_variance, inplace=True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1029 entries, 0 to 1028
Data columns (total 42 columns):
 #   Column                                     Non-Null Count  Dtype
---  -
 0   Age                                       1029 non-null   float64
 1   Attrition                               1029 non-null   int64
 2   Education                               1029 non-null   int64
 3   EnvironmentSatisfaction                 1029 non-null   int64
 4   JobInvolvement                         1029 non-null   int64
 5   JobLevel                               1029 non-null   int64
 6   JobSatisfaction                        1029 non-null   int64
 7   MonthlyIncome                         1029 non-null   int64
 8   MonthlyRate                           1029 non-null   int64
 9   NumCompaniesWorked                    1029 non-null   int64
10   PercentSalaryHike                     1029 non-null   int64
11   PerformanceRating                     1029 non-null   int64
12   RelationshipSatisfaction                1029 non-null   int64
13   StockOptionLevel                      1029 non-null   int64
14   TotalWorkingYears                     1029 non-null   int64
15   TrainingTimesLastYear                 1029 non-null   int64
16   WorkLifeBalance                       1029 non-null   int64
17   YearsAtCompany                        1029 non-null   int64
18   YearsInCurrentRole                    1029 non-null   int64
19   YearsSinceLastPromotion                1029 non-null   int64
20   YearsWithCurrManager                  1029 non-null   int64
21   Gender_Male                           1029 non-null   int64
22   OverTime_Yes                          1029 non-null   int64
23   BusinessTravel_Travel_Frequently       1029 non-null   int64
24   BusinessTravel_Travel_Rarely           1029 non-null   int64
25   Department_Research & Development      1029 non-null   int64
26   Department_Sales                      1029 non-null   int64
27   EducationField_Life Sciences            1029 non-null   int64
28   EducationField_Marketing                1029 non-null   int64
29   EducationField_Medical                 1029 non-null   int64
30   EducationField_Other                   1029 non-null   int64
31   EducationField_Technical Degree         1029 non-null   int64
32   JobRole_Human Resources                 1029 non-null   int64
33   JobRole_Laboratory Technician           1029 non-null   int64
34   JobRole_Manager                       1029 non-null   int64
35   JobRole_Manufacturing Director          1029 non-null   int64
36   JobRole_Research Director              1029 non-null   int64
37   JobRole_Research Scientist              1029 non-null   int64
38   JobRole_Sales Executive                 1029 non-null   int64
39   JobRole_Sales Representative            1029 non-null   int64
40   MaritalStatus_Married                  1029 non-null   int64
41   MaritalStatus_Single                   1029 non-null   int64
dtypes: float64(1), int64(41)
memory usage: 337.8 KB
```

We now have a complete dataset, which is purely numeric, and where we have a moderate set of features where obvious collinearity and lack of variance have been eliminated.

We will next train a logistic regression model and use the coefficients to see if we can further eliminate features.

```
x=df.loc[:, df.columns != 'Attrition'] #feature set
y=df["Attrition"] #target variable

#split the data s.t. we have 20% for testing and 80% for training
X_train, X_test, y_train, y_test = train_test_split(x, y, random_state=2)

#train a lr using liblinear since it is generally good for small training sets
lr = LogisticRegression(solver='liblinear', random_state=2)
lr.fit(X_train, y_train)

print('# of iterations %s' % lr.n_iter_[0])


test_score = lr.score(X_test, y_test)
train_score = lr.score(X_train, y_train)

print('Score on training data: ', train_score)
print('Score on test data: ', test_score)

# of iterations 29
Score on training data: 0.8780804150453956
Score on test data: 0.8643410852713178
```

Now that the model is trained, we can inspect the coefficients.

```
coefficients = lr.coef_
feature_names = X_train.columns
coefficients_df = pd.DataFrame(coefficients, columns=feature_names)
pd.set_option('display.max_columns', None)
coefficients_df.head()
```



	Age	Education	EnvironmentSatisfaction	JobInvolvement	JobLevel	JobSatisfaction	MonthlyIncome	MonthlyRate	NumCompaniesWorked
0	0.00662	-0.004683	-0.311567	-0.290681	-0.057589	-0.220821	-0.000034	0.00002	0.21055

MontlyRate and MonthlyIncome are variables very likely to have collinearty that were missed earlier. We will remove them now.

Education seems to have very low impact on the model. We will therefore remove it in an attempt to simplify the model.

Although Age turns out to seemingly have low impact as well, we will leave it for now as we might could create a new feature by crossing this with some other variable.

```
df.drop(columns=["MonthlyRate", "MonthlyIncome", "Education"], inplace=True)
```

Now we will trian a new model in the same method with the new feature set.

```
x=df.loc[:, df.columns != 'Attrition']
y=df["Attrition"]


X_train, X_test, y_train, y_test = train_test_split(x, y, random_state=2)

lr = LogisticRegression(solver='liblinear', random_state=2)
lr.fit(X_train, y_train)

print('# of iterations %s' % lr.n_iter_[0])

test_score = lr.score(X_test, y_test)
train_score = lr.score(X_train, y_train)

print('Score on training data: ', train_score)
print('Score on test data: ', test_score)
```



```
# of iterations 15
Score on training data: 0.8832684824902723
Score on test data: 0.8914728682170543
```

We have a slight improvement - notably on the test data meaning we have possible reduced overfitting by a bit.

```
#create a dataframe to track our scores on the following models
scores_df=pd.DataFrame(columns=["Model", "Method", "Accuracy"])
```

In the following sections we will take this aproach:

- Use 5-fold cross validation to get the average accuracy for a model
- Use grid search 5 fold cross validation to tune a set of parameters for the model within a restircted value set
- Register the average and best average accuracy for the above in a dataframe

This will be applied across Logistic Regression, Random Forest, and Gradient Boosting models.

```
#use CV on a logistic regression model
cv_results = cross_validate(
    lr, x, y, cv=5,
    scoring=['accuracy', 'precision', 'recall', 'f1'], # Metrics to calculate
    return_train_score=True # If you want to see the train scores as well
)

print("Average accuracy:", cv_results['test_accuracy'].mean())
print("Average precision:", cv_results['test_precision'].mean())
print("Average recall:", cv_results['test_recall'].mean())
print("Average F1 score:", cv_results['test_f1'].mean())
```

```
#save the accuracy
new_row=pd.DataFrame({'Model': ['Logistic Regression'], 'Method': ['Default'], "Accuracy": [ cv_results['test_accuracy'].mean()]})
scores_df = pd.concat([scores_df, new_row], ignore_index=True)

Average accuracy: 0.8717120530428606
Average precision: 0.7335943724044411
Average recall: 0.41460317460317453
Average F1 score: 0.5247012618734046
<ipython-input-355-61170fef1c03>:15: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated.
scores_df = pd.concat([scores_df, new_row], ignore_index=True)
```

```
#use grid search on the logistic regression model
param_grid = {
    'C': [0.01, 0.1, 1, 10],
    'penalty': ['l1', 'l2'],
    'solver': ['liblinear']
}
grid_search = GridSearchCV(lr, param_grid, cv=5, scoring='accuracy')
grid_search.fit(x, y)
print("Best parameters:", grid_search.best_params_)
print("Best cross-validation accuracy:", grid_search.best_score_)

new_row=pd.DataFrame({'Model': ['Logistic Regression'], 'Method': ['Param Search'], "Accuracy": [grid_search.best_score_]})
scores_df = pd.concat([scores_df, new_row], ignore_index=True)
```

```
Best parameters: {'C': 0.1, 'penalty': 'l2', 'solver': 'liblinear'}
Best cross-validation accuracy: 0.8736585365853659
```

```
#use CV on a random forest model
rf=RandomForestClassifier(random_state=42)
cv_results = cross_validate(
    rf, x, y, cv=5,
    scoring=['accuracy', 'precision', 'recall', 'f1'], # Metrics to calculate
    return_train_score=True # If you want to see the train scores as well
)
```

```
print("Average accuracy:", cv_results['test_accuracy'].mean())
print("Average precision:", cv_results['test_precision'].mean())
print("Average recall:", cv_results['test_recall'].mean())
print("Average F1 score:", cv_results['test_f1'].mean())
```

```
new_row=pd.DataFrame({'Model': ['Random Forest'], 'Method': ['Default'], "Accuracy": [cv_results['test_accuracy'].mean()]})
scores_df = pd.concat([scores_df, new_row], ignore_index=True)
```

```
Average accuracy: 0.8503480937721998
Average precision: 0.7758823529411765
Average recall: 0.19857142857142857
Average F1 score: 0.30926819755253276
```

```
#use grid search on a random forest model
param_grid = {
    'n_estimators': [10, 20, 30],
    'max_depth': [5, 10],
    'max_features': ['sqrt', 'log2'],
    'criterion': ['gini', 'entropy']
}
```

```
grid_search = GridSearchCV(rf, param_grid, cv=5, scoring='accuracy')
grid_search.fit(x, y)
print("Best parameters:", grid_search.best_params_)
print("Best cross-validation accuracy:", grid_search.best_score_)
```

```
new_row=pd.DataFrame({'Model': ['Random Forest'], 'Method': ['Param Search'], "Accuracy": [grid_search.best_score_]})
scores_df = pd.concat([scores_df, new_row], ignore_index=True)
```

```
Best parameters: {'criterion': 'entropy', 'max_depth': 10, 'max_features': 'log2', 'n_estimators': 30}
Best cross-validation accuracy: 0.8513189675586077
```

```
#use CV on a gradient boosting model
gb = GradientBoostingClassifier(random_state=42)
cv_results = cross_validate(
    gb, x, y, cv=5,
    scoring=['accuracy', 'precision', 'recall', 'f1'],
    return_train_score=True
```

```

)

print("Average accuracy:", cv_results['test_accuracy'].mean())
print("Average precision:", cv_results['test_precision'].mean())
print("Average recall:", cv_results['test_recall'].mean())
print("Average F1 score:", cv_results['test_f1'].mean())

new_row=pd.DataFrame({'Model': ['Gradient Boosting'], 'Method': ['Default'], "Accuracy": [cv_results['test_accuracy'].mean()]})
scores_df = pd.concat([scores_df, new_row], ignore_index=True)

➞ Average accuracy: 0.8658726024153445
Average precision: 0.7259006353124
Average recall: 0.3685714285714286
Average F1 score: 0.4807791811247844

#use grid search on a gradient boosting model
param_grid = {
    'n_estimators': [10, 20, 30],
    'max_depth': [5, 10],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1.0],
    'max_features': ['sqrt', 'log2']
}

grid_search = GridSearchCV(gb, param_grid, cv=5, scoring='accuracy')
grid_search.fit(x, y)
print("Best parameters:", grid_search.best_params_)
print("Best cross-validation accuracy:", grid_search.best_score_)

new_row=pd.DataFrame({'Model': ['Gradient Boosting'], 'Method': ['Param Search'], "Accuracy": [grid_search.best_score_]})
scores_df = pd.concat([scores_df, new_row], ignore_index=True)

➞ Best parameters: {'learning_rate': 0.2, 'max_depth': 5, 'max_features': 'log2', 'n_estimators': 30, 'subsample': 1.0}
Best cross-validation accuracy: 0.8590954297892492

```

Let's take a look at the results, graphically.

```


plt.figure(figsize=(10, 6))
sns.barplot(data=scores_df, x='Model', y='Accuracy', hue='Method')

plt.title('Model Accuracy Comparison')
plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.ylim(0.8, 0.9)
plt.legend(title='Method')
plt.grid(axis='y')

plt.show()

```

```
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:949: FutureWarning: When grouping with a length-1 list-like, you will need to p
data_subset = grouped_data.get_group(pd_key)
/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:949: FutureWarning: When grouping with a length-1 list-like, you will need to p
data_subset = grouped_data.get_group(pd_key)
```



The overall best model was a parameter-tuned Liblinear Logistic Regression classifier.